

The Red Hat newlib C Math Library

libm 1.16.0
December 2007

Steve Chamberlain
Roland Pesch
Red Hat Support
Jeff Johnston

Red Hat Support
sac@cygnus.com
pesch@cygnus.com
jjohnstn@redhat.com

Copyright © 1992, 1993, 1994-2004 Red Hat, Inc.

‘libm’ includes software developed at SunPro, a Sun Microsystems, Inc. business. Permission to use, copy, modify, and distribute this software is freely granted, provided that this notice is preserved.

Permission is granted to make and distribute verbatim copies of this manual provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this manual under the conditions for verbatim copying, subject to the terms of the GNU General Public License, which includes the provision that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this manual into another language, under the above conditions for modified versions.

1 Mathematical Functions (`'math.h'`)

This chapter groups a wide variety of mathematical functions. The corresponding definitions and declarations are in `'math.h'`. Two definitions from `'math.h'` are of particular interest.

1. The representation of infinity as a `double` is defined as `HUGE_VAL`; this number is returned on overflow by many functions.
2. The structure `exception` is used when you write customized error handlers for the mathematical functions. You can customize error handling for most of these functions by defining your own version of `matherr`; see the section on `matherr` for details.

Since the error handling code calls `fputs`, the mathematical subroutines require stubs or minimal implementations for the same list of OS subroutines as `fputs`: `close`, `fstat`, `isatty`, `lseek`, `read`, `sbrk`, `write`. See [section “System Calls” in *The Red Hat newlib C Library*](#), for a discussion and for sample minimal implementations of these support subroutines.

Alternative declarations of the mathematical functions, which exploit specific machine capabilities to operate faster—but generally have less error checking and may reflect additional limitations on some machines—are available when you include `'fastmath.h'` instead of `'math.h'`.

1.1 Version of library

There are four different versions of the math library routines: IEEE, POSIX, X/Open, or SVID. The version may be selected at runtime by setting the global variable `_LIB_VERSION`, defined in `'math.h'`. It may be set to one of the following constants defined in `'math.h'`: `_IEEE_`, `_POSIX_`, `_XOPEN_`, or `_SVID_`. The `_LIB_VERSION` variable is not specific to any thread, and changing it will affect all threads.

The versions of the library differ only in how errors are handled.

In IEEE mode, the `matherr` function is never called, no warning messages are printed, and `errno` is never set.

In POSIX mode, `errno` is set correctly, but the `matherr` function is never called and no warning messages are printed.

In X/Open mode, `errno` is set correctly, and `matherr` is called, but warning message are not printed.

In SVID mode, functions which overflow return 3.40282346638528860e+38, the maximum single-precision floating-point value, rather than infinity. Also, `errno` is set correctly, `matherr` is called, and, if `matherr` returns 0, warning messages are printed for some errors. For example, by default `'log(-1.0)'` writes this message on standard error output:

`log: DOMAIN error`

The library is set to X/Open mode by default.

1.2 `acos`, `acosf`—arc cosine

Synopsis

```
#include <math.h>
double acos(double x);
float acosf(float x);
```

Description

`acos` computes the inverse cosine (arc cosine) of the input value. Arguments to `acos` must be in the range -1 to 1 .

`acosf` is identical to `acos`, except that it performs its calculations on `floats`.

Returns

`acos` and `acosf` return values in radians, in the range of 0 to π .

If x is not between -1 and 1 , the returned value is NaN (not a number) the global variable `errno` is set to `EDOM`, and a **DOMAIN error** message is sent as standard error output.

You can modify error handling for these functions using `matherr`.

1.3 `acosh`, `acoshf`—inverse hyperbolic cosine

Synopsis

```
#include <math.h>
double acosh(double x);
float acoshf(float x);
```

Description

`acosh` calculates the inverse hyperbolic cosine of `x`. `acosh` is defined as

$$\ln\left(x + \sqrt{x^2 - 1}\right)$$

`x` must be a number greater than or equal to 1.

`acoshf` is identical, other than taking and returning floats.

Returns

`acosh` and `acoshf` return the calculated value. If `x` less than 1, the return value is NaN and `errno` is set to `EDOM`.

You can change the error-handling behavior with the non-ANSI `matherr` function.

Portability

Neither `acosh` nor `acoshf` are ANSI C. They are not recommended for portable programs.

1.4 `asin`, `asinf`—arc sine

Synopsis

```
#include <math.h>
double asin(double x);
float asinf(float x);
```

Description

`asin` computes the inverse sine (arc sine) of the argument `x`. Arguments to `asin` must be in the range -1 to 1 .

`asinf` is identical to `asin`, other than taking and returning floats.

You can modify error handling for these routines using `matherr`.

Returns

`asin` returns values in radians, in the range of $-\pi/2$ to $\pi/2$.

If `x` is not in the range -1 to 1 , `asin` and `asinf` return NaN (not a number), set the global variable `errno` to `EDOM`, and issue a `DOMAIN error` message.

You can change this error treatment using `matherr`.

1.5 asinh, asinhf—inverse hyperbolic sine

Synopsis

```
#include <math.h>
double asinh(double x);
float asinhf(float x);
```

Description

asinh calculates the inverse hyperbolic sine of *x*. **asinh** is defined as

$$\text{sign}(x) \times \ln\left(|x| + \sqrt{1 + x^2}\right)$$

asinhf is identical, other than taking and returning floats.

Returns

asinh and **asinhf** return the calculated value.

Portability

Neither **asinh** nor **asinhf** are ANSI C.

1.6 `atan`, `atanf`—arc tangent

Synopsis

```
#include <math.h>
double atan(double x);
float atanf(float x);
```

Description

`atan` computes the inverse tangent (arc tangent) of the input value.

`atanf` is identical to `atan`, save that it operates on `floats`.

Returns

`atan` returns a value in radians, in the range of $-\pi/2$ to $\pi/2$.

Portability

`atan` is ANSI C. `atanf` is an extension.

1.7 atan2, atan2f—arc tangent of y/x

Synopsis

```
#include <math.h>
double atan2(double y,double x);
float atan2f(float y,float x);
```

Description

atan2 computes the inverse tangent (arc tangent) of y/x . **atan2** produces the correct result even for angles near $\pi/2$ or $-\pi/2$ (that is, when x is near 0).

atan2f is identical to **atan2**, save that it takes and returns **float**.

Returns

atan2 and **atan2f** return a value in radians, in the range of $-\pi$ to π .

You can modify error handling for these functions using **matherr**.

Portability

atan2 is ANSI C. **atan2f** is an extension.

1.8 atanh, atanhf—inverse hyperbolic tangent

Synopsis

```
#include <math.h>
double atanh(double x);
float atanhf(float x);
```

Description

atanh calculates the inverse hyperbolic tangent of x .

atanhf is identical, other than taking and returning **float** values.

Returns

atanh and **atanhf** return the calculated value.

If $|x|$ is greater than 1, the global **errno** is set to **EDOM** and the result is a NaN. A **DOMAIN error** is reported.

If $|x|$ is 1, the global **errno** is set to **EDOM**; and the result is infinity with the same sign as x . A **SING error** is reported.

You can modify the error handling for these routines using **matherr**.

Portability

Neither **atanh** nor **atanhf** are ANSI C.

1.9 jN, jNf, yN, yNf—Bessel functions

Synopsis

```
#include <math.h>
double j0(double x);
float j0f(float x);
double j1(double x);
float j1f(float x);
double jn(int n, double x);
float jnf(int n, float x);
double y0(double x);
float y0f(float x);
double y1(double x);
float y1f(float x);
double yn(int n, double x);
float ynf(int n, float x);
```

Description

The Bessel functions are a family of functions that solve the differential equation

$$x^2 \frac{d^2 y}{dx^2} + x \frac{dy}{dx} + (x^2 - p^2)y = 0$$

These functions have many applications in engineering and physics.

`jn` calculates the Bessel function of the first kind of order n . `j0` and `j1` are special cases for order 0 and order 1 respectively.

Similarly, `yn` calculates the Bessel function of the second kind of order n , and `y0` and `y1` are special cases for order 0 and 1.

`jnf`, `j0f`, `j1f`, `ynf`, `y0f`, and `y1f` perform the same calculations, but on `float` rather than `double` values.

Returns

The value of each Bessel function at x is returned.

Portability

None of the Bessel functions are in ANSI C.

1.10 `cbrt`, `cbrtf`—cube root

Synopsis

```
#include <math.h>
double cbrt(double x);
float  cbrtf(float x);
```

Description

`cbrt` computes the cube root of the argument.

Returns

The cube root is returned.

Portability

`cbrt` is in System V release 4. `cbrtf` is an extension.

1.11 copysign, copysignf—sign of *y*, magnitude of *x*

Synopsis

```
#include <math.h>
double copysign (double x, double y);
float copysignf (float x, float y);
```

Description

`copysign` constructs a number with the magnitude (absolute value) of its first argument, *x*, and the sign of its second argument, *y*.

`copysignf` does the same thing; the two functions differ only in the type of their arguments and result.

Returns

`copysign` returns a `double` with the magnitude of *x* and the sign of *y*. `copysignf` returns a `float` with the magnitude of *x* and the sign of *y*.

Portability

`copysign` is not required by either ANSI C or the System V Interface Definition (Issue 2).

1.12 `cosh`, `coshf`—hyperbolic cosine

Synopsis

```
#include <math.h>
double cosh(double x);
float coshf(float x)
```

Description

`cosh` computes the hyperbolic cosine of the argument `x`. `cosh(x)` is defined as

$$\frac{(e^x + e^{-x})}{2}$$

Angles are specified in radians. `coshf` is identical, save that it takes and returns `float`.

Returns

The computed value is returned. When the correct value would create an overflow, `cosh` returns the value `HUGE_VAL` with the appropriate sign, and the global value `errno` is set to `ERANGE`.

You can modify error handling for these functions using the function `matherr`.

Portability

`cosh` is ANSI. `coshf` is an extension.

1.13 erf, erff, erfc, erfcf—error function

Synopsis

```
#include <math.h>
double erf(double x);
float erff(float x);
double erfc(double x);
float erfcf(float x);
```

Description

erf calculates an approximation to the “error function”, which estimates the probability that an observation will fall within *x* standard deviations of the mean (assuming a normal distribution). The error function is defined as

$$\frac{2}{\sqrt{\pi}} \times \int_0^x e^{-t^2} dt$$

erfc calculates the complementary probability; that is, **erfc(x)** is **1 - erf(x)**. **erfc** is computed directly, so that you can use it to avoid the loss of precision that would result from subtracting large probabilities (on large *x*) from 1.

erff and **erfcf** differ from **erf** and **erfc** only in the argument and result types.

Returns

For positive arguments, **erf** and all its variants return a probability—a number between 0 and 1.

Portability

None of the variants of **erf** are ANSI C.

1.14 exp, expf—exponential

Synopsis

```
#include <math.h>
double exp(double x);
float expf(float x);
```

Description

`exp` and `expf` calculate the exponential of x , that is, e^x (where e is the base of the natural system of logarithms, approximately 2.71828).

You can use the (non-ANSI) function `matherr` to specify error handling for these functions.

Returns

On success, `exp` and `expf` return the calculated value. If the result underflows, the returned value is 0. If the result overflows, the returned value is `HUGE_VAL`. In either case, `errno` is set to `ERANGE`.

Portability

`exp` is ANSI C. `expf` is an extension.

1.15 `expm1`, `expm1f`—exponential minus 1

Synopsis

```
#include <math.h>
double expm1(double x);
float expm1f(float x);
```

Description

`expm1` and `expm1f` calculate the exponential of x and subtract 1, that is, $e^x - 1$ (where e is the base of the natural system of logarithms, approximately 2.71828). The result is accurate even for small values of x , where using `exp(x)-1` would lose many significant digits.

Returns

e raised to the power x , minus 1.

Portability

Neither `expm1` nor `expm1f` is required by ANSI C or by the System V Interface Definition (Issue 2).

1.16 `fabs`, `fabsf`—absolute value (magnitude)

Synopsis

```
#include <math.h>
double fabs(double x);
float fabsf(float x);
```

Description

`fabs` and `fabsf` calculate $|x|$, the absolute value (magnitude) of the argument `x`, by direct manipulation of the bit representation of `x`.

Returns

The calculated value is returned. No errors are detected.

Portability

`fabs` is ANSI. `fabsf` is an extension.

1.17 floor, floorf, ceil, ceilf—floor and ceiling

Synopsis

```
#include <math.h>
double floor(double x);
float floorf(float x);
double ceil(double x);
float ceilf(float x);
```

Description

`floor` and `floorf` find $\lfloor x \rfloor$, the nearest integer less than or equal to x . `ceil` and `ceilf` find $\lceil x \rceil$, the nearest integer greater than or equal to x .

Returns

`floor` and `ceil` return the integer result as a double. `floorf` and `ceilf` return the integer result as a float.

Portability

`floor` and `ceil` are ANSI. `floorf` and `ceilf` are extensions.

1.18 `fmod`, `fmodf`—floating-point remainder (modulo)

Synopsis

```
#include <math.h>
double fmod(double x, double y)
float fmodf(float x, float y)
```

Description

The `fmod` and `fmodf` functions compute the floating-point remainder of x/y (x modulo y).

Returns

The `fmod` function returns the value $x - i \times y$, for the largest integer i such that, if y is nonzero, the result has the same sign as x and magnitude less than the magnitude of y .

`fmod(x,0)` returns NaN, and sets `errno` to `EDOM`.

You can modify error treatment for these functions using `matherr`.

Portability

`fmod` is ANSI C. `fmodf` is an extension.

1.19 frexp, frexpf—split floating-point number

Synopsis

```
#include <math.h>
double frexp(double val, int *exp);
float frexpf(float val, int *exp);
```

Description

All nonzero, normal numbers can be described as $m * 2^p$. **frexp** represents the double *val* as a mantissa *m* and a power of two *p*. The resulting mantissa will always be greater than or equal to 0.5, and less than 1.0 (as long as *val* is nonzero). The power of two will be stored in **exp*.

m and *p* are calculated so that $val = m \times 2^p$.

frexpf is identical, other than taking and returning floats rather than doubles.

Returns

frexp returns the mantissa *m*. If *val* is 0, infinity, or Nan, **frexp** will set **exp* to 0 and return *val*.

Portability

frexp is ANSI. **frexpf** is an extension.

1.20 gamma, gammaf, lgamma, lgammaf, gamma_r,

Synopsis

```
#include <math.h>
double gamma(double x);
float gammaf(float x);
double lgamma(double x);
float lgammaf(float x);
double gamma_r(double x, int *signgam);
float gammaf_r(float x, int *signgam);
double lgamma_r(double x, int *signgam);
float lgammaf_r(float x, int *signgam);
```

Description

gamma calculates $\ln(\Gamma(x))$, the natural logarithm of the gamma function of x . The gamma function (`exp(gamma(x))`) is a generalization of factorial, and retains the property that $\Gamma(N) \equiv N \times \Gamma(N - 1)$. Accordingly, the results of the gamma function itself grow very quickly. **gamma** is defined as $\ln(\Gamma(x))$ rather than simply $\Gamma(x)$ to extend the useful range of results representable.

The sign of the result is returned in the global variable **signgam**, which is declared in **math.h**.

gammaf performs the same calculation as **gamma**, but uses and returns **float** values.

lgamma and **lgammaf** are alternate names for **gamma** and **gammaf**. The use of **lgamma** instead of **gamma** is a reminder that these functions compute the log of the gamma function, rather than the gamma function itself.

The functions **gamma_r**, **gammaf_r**, **lgamma_r**, and **lgammaf_r** are just like **gamma**, **gammaf**, **lgamma**, and **lgammaf**, respectively, but take an additional argument. This additional argument is a pointer to an integer. This additional argument is used to return the sign of the result, and the global variable **signgam** is not used. These functions may be used for reentrant calls (but they will still set the global variable **errno** if an error occurs).

Returns

Normally, the computed result is returned.

When x is a nonpositive integer, **gamma** returns **HUGE_VAL** and **errno** is set to **EDOM**. If the result overflows, **gamma** returns **HUGE_VAL** and **errno** is set to **ERANGE**.

You can modify this error treatment using **matherr**.

Portability

Neither **gamma** nor **gammaf** is ANSI C.

1.21 hypot, hypotf—distance from origin

Synopsis

```
#include <math.h>
double hypot(double x, double y);
float hypotf(float x, float y);
```

Description

`hypot` calculates the Euclidean distance $\sqrt{x^2 + y^2}$ between the origin (0,0) and a point represented by the Cartesian coordinates (x,y). `hypotf` differs only in the type of its arguments and result.

Returns

Normally, the distance value is returned. On overflow, `hypot` returns `HUGE_VAL` and sets `errno` to `ERANGE`.

You can change the error treatment with `matherr`.

Portability

`hypot` and `hypotf` are not ANSI C.

1.22 `ilogb`, `ilogbf`—get exponent of floating-point number

Synopsis

```
#include <math.h>
int ilogb(double val);
int ilogbf(float val);
```

Description

All nonzero, normal numbers can be described as $m * 2^{**}p$. `ilogb` and `ilogbf` examine the argument *val*, and return *p*. The functions `frexp` and `frexpf` are similar to `ilogb` and `ilogbf`, but also return *m*.

Returns

`ilogb` and `ilogbf` return the power of two used to form the floating-point argument. If *val* is 0, they return `- INT_MAX` (`INT_MAX` is defined in `limits.h`). If *val* is infinite, or NaN, they return `INT_MAX`.

Portability

Neither `ilogb` nor `ilogbf` is required by ANSI C or by the System V Interface Definition (Issue 2).

1.23 `infinity`, `infinityf`—representation of infinity

Synopsis

```
#include <math.h>
double infinity(void);
float infinityf(void);
```

Description

`infinity` and `infinityf` return the special number IEEE infinity in double- and single-precision arithmetic respectively.

1.24 isnan, isnanf, isinf, isinff, finite, finitelf—test for exceptional numbers

Synopsis

```
#include <ieeefp.h>
int isnan(double arg);
int isinf(double arg);
int finite(double arg);
int isnanf(float arg);
int isinff(float arg);
int finitelf(float arg);
```

Description

These functions provide information on the floating-point argument supplied.

There are five major number formats:

zero A number which contains all zero bits.

subnormal A number with a zero exponent but a nonzero fraction.

normal A number with an exponent and a fraction.

infinity A number with an all 1's exponent and a zero fraction.

NAN A number with an all 1's exponent and a nonzero fraction.

isnan returns 1 if the argument is a nan. **isinf** returns 1 if the argument is infinity. **finite** returns 1 if the argument is zero, subnormal or normal. The **isnanf**, **isinff** and **finitelf** functions perform the same operations as their **isnan**, **isinf** and **finite** counterparts, but on single-precision floating-point numbers.

It should be noted that the C99 standard dictates that **isnan** and **isinf** are macros that operate on multiple types of floating-point. The SUSv2 standard declares **isnan** as a function taking double. Newlib has decided to declare them both as macros in **math.h** and as functions in **ieeefp.h**.

1.25 ldexp, ldexpf—load exponent

Synopsis

```
#include <math.h>
double ldexp(double val, int exp);
float ldexpf(float val, int exp);
```

Description

`ldexp` calculates the value $val \times 2^{exp}$. `ldexpf` is identical, save that it takes and returns `float` rather than `double` values.

Returns

`ldexp` returns the calculated value.

Underflow and overflow both set `errno` to `ERANGE`. On underflow, `ldexp` and `ldexpf` return 0.0. On overflow, `ldexp` returns plus or minus `HUGE_VAL`.

Portability

`ldexp` is ANSI. `ldexpf` is an extension.

1.26 `log`, `logf`—natural logarithms

Synopsis

```
#include <math.h>
double log(double x);
float logf(float x);
```

Description

Return the natural logarithm of x , that is, its logarithm base e (where e is the base of the natural system of logarithms, 2.71828...). `log` and `logf` are identical save for the return and argument types.

You can use the (non-ANSI) function `matherr` to specify error handling for these functions.

Returns

Normally, returns the calculated value. When x is zero, the returned value is `-HUGE_VAL` and `errno` is set to `ERANGE`. When x is negative, the returned value is NaN (not a number) and `errno` is set to `EDOM`. You can control the error behavior via `matherr`.

Portability

`log` is ANSI. `logf` is an extension.

1.27 `log10`, `log10f`—base 10 logarithms

Synopsis

```
#include <math.h>
double log10(double x);
float log10f(float x);
```

Description

`log10` returns the base 10 logarithm of `x`. It is implemented as $\log(x) / \log(10)$.

`log10f` is identical, save that it takes and returns `float` values.

Returns

`log10` and `log10f` return the calculated value.

See the description of `log` for information on errors.

Portability

`log10` is ANSI C. `log10f` is an extension.

1.28 `log1p`, `log1pf`—log of $1 + x$

Synopsis

```
#include <math.h>
double log1p(double x);
float log1pf(float x);
```

Description

`log1p` calculates $\ln(1 + x)$, the natural logarithm of $1+x$. You can use `log1p` rather than `'log(1+x)'` for greater precision when x is very small.

`log1pf` calculates the same thing, but accepts and returns `float` values rather than `double`.

Returns

`log1p` returns a `double`, the natural log of $1+x$. `log1pf` returns a `float`, the natural log of $1+x$.

Portability

Neither `log1p` nor `log1pf` is required by ANSI C or by the System V Interface Definition (Issue 2).

1.29 matherr—modifiable math error handler

Synopsis

```
#include <math.h>
int matherr(struct exception *e);
```

Description

matherr is called whenever a math library function generates an error. You can replace **matherr** by your own subroutine to customize error treatment. The customized **matherr** must return 0 if it fails to resolve the error, and non-zero if the error is resolved.

When **matherr** returns a nonzero value, no error message is printed and the value of **errno** is not modified. You can accomplish either or both of these things in your own **matherr** using the information passed in the structure ***e**.

This is the **exception** structure (defined in ‘**math.h**’):

```
struct exception {
    int type;
    char *name;
    double arg1, arg2, retval;
    int err;
};
```

The members of the exception structure have the following meanings:

- | | |
|-------------------|--|
| type | The type of mathematical error that occurred; macros encoding error types are also defined in ‘ math.h ’. |
| name | a pointer to a null-terminated string holding the name of the math library function where the error occurred. |
| arg1, arg2 | The arguments which caused the error. |
| retval | The error return value (what the calling function will return). |
| err | If set to be non-zero, this is the new value assigned to errno . |

The error types defined in ‘**math.h**’ represent possible mathematical errors as follows:

- | | |
|------------------|--|
| DOMAIN | An argument was not in the domain of the function; e.g. log(-1.0) . |
| SING | The requested calculation would result in a singularity; e.g. pow(0.0,-2.0) . |
| OVERFLOW | A calculation would produce a result too large to represent; e.g. exp(1000.0) . |
| UNDERFLOW | A calculation would produce a result too small to represent; e.g. exp(-1000.0) . |
| TLOSS | Total loss of precision. The result would have no significant digits; e.g. sin(10e70) . |
| PLOSS | Partial loss of precision. |

Returns

The library definition for **matherr** returns 0 in all cases.

You can change the calling function's result from a customized `matherr` by modifying `e->retval`, which propagates back to the caller.

If `matherr` returns 0 (indicating that it was not able to resolve the error) the caller sets `errno` to an appropriate value, and prints an error message.

Portability

`matherr` is not ANSI C.

1.30 `modf`, `modff`—split fractional and integer parts

Synopsis

```
#include <math.h>
double modf(double val, double *ipart);
float modff(float val, float *ipart);
```

Description

`modf` splits the double `val` apart into an integer part and a fractional part, returning the fractional part and storing the integer part in `*ipart`. No rounding whatsoever is done; the sum of the integer and fractional parts is guaranteed to be exactly equal to `val`. That is, if `realpart = modf(val, &intpart)`; then ‘`realpart+intpart`’ is the same as `val`. `modff` is identical, save that it takes and returns `float` rather than `double` values.

Returns

The fractional part is returned. Each result has the same sign as the supplied argument `val`.

Portability

`modf` is ANSI C. `modff` is an extension.

1.31 `nan`, `nanf`—representation of “Not a Number”

Synopsis

```
#include <math.h>
double nan(const char *);
float nanf(const char *);
```

Description

`nan` and `nanf` return an IEEE NaN (Not a Number) in double- and single-precision arithmetic respectively. The argument is currently disregarded.

1.32 nextafter, nextafterf—get next number

Synopsis

```
#include <math.h>
double nextafter(double val, double dir);
float nextafterf(float val, float dir);
```

Description

nextafter returns the double-precision floating-point number closest to *val* in the direction toward *dir*. **nextafterf** performs the same operation in single precision. For example, **nextafter**(0.0,1.0) returns the smallest positive number which is representable in double precision.

Returns

Returns the next closest number to *val* in the direction toward *dir*.

Portability

Neither **nextafter** nor **nextafterf** is required by ANSI C or by the System V Interface Definition (Issue 2).

1.33 `pow`, `powf`— x to the power y

Synopsis

```
#include <math.h>
double pow(double x, double y);
float powf(float x, float y);
```

Description

`pow` and `powf` calculate x raised to the exponent y . (That is, x^y .)

Returns

On success, `pow` and `powf` return the value calculated.

When the argument values would produce overflow, `pow` returns `HUGE_VAL` and set `errno` to `ERANGE`. If the argument x passed to `pow` or `powf` is a negative noninteger, and y is also not an integer, then `errno` is set to `EDOM`. If x and y are both 0, then `pow` and `powf` return 1.

You can modify error handling for these functions using `matherr`.

Portability

`pow` is ANSI C. `powf` is an extension.

1.34 remainder, remainderf—round and remainder

Synopsis

```
#include <math.h>
double remainder(double x, double y);
float remainderf(float x, float y);
```

Description

`remainder` and `remainderf` find the remainder of x/y ; this value is in the range $-y/2$.. $+y/2$.

Returns

`remainder` returns the integer result as a double.

Portability

`remainder` is a System V release 4. `remainderf` is an extension.

1.35 `scalbn`, `scalbnf`—scale by power of two

Synopsis

```
#include <math.h>
double scalbn(double x, int y);
float scalbnf(float x, int y);
```

Description

`scalbn` and `scalbnf` scale x by n , returning x times 2 to the power n . The result is computed by manipulating the exponent, rather than by actually performing an exponentiation or multiplication.

Returns

x times 2 to the power n .

Portability

Neither `scalbn` nor `scalbnf` is required by ANSI C or by the System V Interface Definition (Issue 2).

1.36 `sqrt`, `sqrtf`—positive square root

Synopsis

```
#include <math.h>
double sqrt(double x);
float sqrtf(float x);
```

Description

`sqrt` computes the positive square root of the argument. You can modify error handling for this function with `matherr`.

Returns

On success, the square root is returned. If `x` is real and positive, then the result is positive. If `x` is real and negative, the global value `errno` is set to `EDOM` (domain error).

Portability

`sqrt` is ANSI C. `sqrtf` is an extension.

1.37 sin, sinf, cos, cosf—sine or cosine

Synopsis

```
#include <math.h>
double sin(double x);
float  sinf(float x);
double cos(double x);
float  cosf(float x);
```

Description

`sin` and `cos` compute (respectively) the sine and cosine of the argument `x`. Angles are specified in radians.

`sinf` and `cosf` are identical, save that they take and return `float` values.

Returns

The sine or cosine of `x` is returned.

Portability

`sin` and `cos` are ANSI C. `sinf` and `cosf` are extensions.

1.38 `sinh`, `sinhf`—hyperbolic sine

Synopsis

```
#include <math.h>
double sinh(double x);
float  sinhf(float x);
```

Description

`sinh` computes the hyperbolic sine of the argument `x`. Angles are specified in radians. `sinh(x)` is defined as

$$\frac{e^x - e^{-x}}{2}$$

`sinhf` is identical, save that it takes and returns `float` values.

Returns

The hyperbolic sine of `x` is returned.

When the correct result is too large to be representable (an overflow), `sinh` returns `HUGE_VAL` with the appropriate sign, and sets the global value `errno` to `ERANGE`.

You can modify error handling for these functions with `matherr`.

Portability

`sinh` is ANSI C. `sinhf` is an extension.

1.39 `tan`, `tanf`—tangent

Synopsis

```
#include <math.h>
double tan(double x);
float tanf(float x);
```

Description

`tan` computes the tangent of the argument `x`. Angles are specified in radians.

`tanf` is identical, save that it takes and returns `float` values.

Returns

The tangent of `x` is returned.

Portability

`tan` is ANSI. `tanf` is an extension.

1.40 `tanh`, `tanhf`—hyperbolic tangent

Synopsis

```
#include <math.h>
double tanh(double x);
float tanhf(float x);
```

Description

`tanh` computes the hyperbolic tangent of the argument `x`. Angles are specified in radians.

`tanh(x)` is defined as

$$\sinh(x)/\cosh(x)$$

`tanhf` is identical, save that it takes and returns `float` values.

Returns

The hyperbolic tangent of `x` is returned.

Portability

`tanh` is ANSI C. `tanhf` is an extension.

2 Reentrancy Properties of `libm`

When a `libm` function detects an exceptional case, `errno` may be set, the `matherr` function may be called, and a error message may be written to the standard error stream. This behavior may not be reentrant.

With reentrant C libraries like the Red Hat newlib C library, `errno` is a macro which expands to the per-thread error value. This makes it thread safe.

When the user provides his own `matherr` function it must be reentrant for the math library as a whole to be reentrant.

In normal debugged programs, there are usually no math subroutine errors—and therefore no assignments to `errno` and no `matherr` calls; in that situation, the math functions behave reentrantly.

Library Index

A

acos	3
acosf	3
acosh	4
acoshf	4
asin	5
asinf	5
asinh	6
asinhf	6
atan	7
atan2	8
atan2f	8
atanf	7
atanh	9
atanhf	9

C

cbrt	11
cbrtf	11
ceil	18
ceilf	18
copysign	12
copysignf	12
cos	39
cosf	39

E

erf	14
erfc	14
erfcf	14
erff	14
exp	15
expf	15
expm1	16
expm1f	16

F

fabs	17
fabsf	17
finite	25
finitef	25
floor	18
floorf	18
fmod	19
fmodf	19
frexp	20
frexpf	20

G

gamma	21
-------	----

gamma_r	21
gammaf	21
gammaf_r	21

H

hypot	22
hypotf	22

I

ilogb	23
ilogbf	23
infinity	24
infinityf	24
isinf	25
isinff	25
isnan	25
isnanf	25

J

j0	10
j0f	10
j1	10
j1f	10
jn	10
jnf	10

L

ldexp	26
ldexpf	26
lgamma	21
lgamma_r	21
lgammaf	21
lgammaf_r	21
log	27
log10	28
log10f	28
log1p	29
log1pf	29
logf	27

M

matherr	30
matherr and reentrancy	43
modf	32
modff	32

N

nan	33
-----	----

<code>nanf</code>	33
<code>nextafter</code>	34
<code>nextafterf</code>	34

O

OS stubs	1
----------------	---

P

<code>pow</code>	35
<code>powf</code>	35

R

reentrancy	43
<code>remainder</code>	36
<code>remainderf</code>	36

S

<code>scalbn</code>	37
<code>scalbnf</code>	37
<code>sin</code>	39

<code>sinf</code>	39
<code>sinh</code>	40
<code>sinhf</code>	40
<code>sqrt</code>	38
<code>sqrtf</code>	38
stubs	1
support subroutines	1
system calls	1

T

<code>tan</code>	41
<code>tanf</code>	41
<code>tanh</code>	42
<code>tanhf</code>	42

Y

<code>y0</code>	10
<code>y0f</code>	10
<code>y1</code>	10
<code>y1f</code>	10
<code>yn</code>	10
<code>ynf</code>	10

The body of this manual is set in
cmr10 at 10.95pt,
with headings in **cmb10 at 10.95pt**
and examples in cmr10 at 10.95pt.
cmr10 at 10.95pt and
cml10 at 10.95pt
are used for emphasis.

Table of Contents

1	Mathematical Functions (<code>'math.h'</code>)	1
1.1	Version of library	2
1.2	<code>acos</code> , <code>acosf</code> —arc cosine	3
1.3	<code>acosh</code> , <code>acoshf</code> —inverse hyperbolic cosine	4
1.4	<code>asin</code> , <code>asinf</code> —arc sine	5
1.5	<code>asinh</code> , <code>asinhf</code> —inverse hyperbolic sine	6
1.6	<code>atan</code> , <code>atanf</code> —arc tangent	7
1.7	<code>atan2</code> , <code>atan2f</code> —arc tangent of y/x	8
1.8	<code>atanh</code> , <code>atanhf</code> —inverse hyperbolic tangent	9
1.9	<code>jN</code> , <code>jNf</code> , <code>yN</code> , <code>yNf</code> —Bessel functions	10
1.10	<code>cbrt</code> , <code>cbrtf</code> —cube root	11
1.11	<code>copysign</code> , <code>copysignf</code> —sign of y , magnitude of x	12
1.12	<code>cosh</code> , <code>coshf</code> —hyperbolic cosine	13
1.13	<code>erf</code> , <code>erff</code> , <code>erfc</code> , <code>erfcf</code> —error function	14
1.14	<code>exp</code> , <code>expf</code> —exponential	15
1.15	<code>expm1</code> , <code>expm1f</code> —exponential minus 1	16
1.16	<code>fabs</code> , <code>fabsf</code> —absolute value (magnitude)	17
1.17	<code>floor</code> , <code>floorf</code> , <code>ceil</code> , <code>ceilf</code> —floor and ceiling	18
1.18	<code>fmod</code> , <code>fmodf</code> —floating-point remainder (modulo)	19
1.19	<code>frexp</code> , <code>frexpf</code> —split floating-point number	20
1.20	<code>gamma</code> , <code>gammaf</code> , <code>lgamma</code> , <code>lgammaf</code> , <code>gamma_r</code> ,	21
1.21	<code>hypot</code> , <code>hypotf</code> —distance from origin	22
1.22	<code>ilogb</code> , <code>ilogbf</code> —get exponent of floating-point number	23
1.23	<code>infinity</code> , <code>infinityf</code> —representation of infinity	24
1.24	<code>isnan</code> , <code>isnanf</code> , <code>isinf</code> , <code>isinff</code> , <code>finite</code> , <code>finitef</code> —test for exceptional numbers	25
1.25	<code>ldexp</code> , <code>ldexpf</code> —load exponent	26
1.26	<code>log</code> , <code>logf</code> —natural logarithms	27
1.27	<code>log10</code> , <code>log10f</code> —base 10 logarithms	28
1.28	<code>log1p</code> , <code>log1pf</code> —log of $1 + x$	29
1.29	<code>matherr</code> —modifiable math error handler	30
1.30	<code>modf</code> , <code>modff</code> —split fractional and integer parts	32
1.31	<code>nan</code> , <code>nanf</code> —representation of “Not a Number”	33
1.32	<code>nextafter</code> , <code>nextafterf</code> —get next number	34
1.33	<code>pow</code> , <code>powf</code> — x to the power y	35
1.34	<code>remainder</code> , <code>remainderf</code> —round and remainder	36
1.35	<code>scalbn</code> , <code>scalbnf</code> —scale by power of two	37
1.36	<code>sqrt</code> , <code>sqrtf</code> —positive square root	38
1.37	<code>sin</code> , <code>sinf</code> , <code>cos</code> , <code>cosf</code> —sine or cosine	39
1.38	<code>sinh</code> , <code>sinhf</code> —hyperbolic sine	40
1.39	<code>tan</code> , <code>tanf</code> —tangent	41
1.40	<code>tanh</code> , <code>tanhf</code> —hyperbolic tangent	42

2	Reentrancy Properties of libm	43
----------	--	-----------

	Library Index	45
--	----------------------------	-----------